

Original Mixed-radix: Final Evidence and Closure

A real mechanism, but no stable natural-data method

2026-08-04

1. The Answer First

Final result for the tested method

The original fixed-adjacent mixed-radix method is **not competitive on SIFT1M or GIST1M** under the frozen Recall–speed evaluation.

- At 32 bytes per database vector, it becomes exactly the power-of-two control.
- At 64 bytes per database vector, Recall changes by at most about one tenth of one percentage point, with both positive and negative signs.
- A synthetic positive control now confirms that the intended mechanism itself is real and reaches the query result.

Interpretation

The implementation can exploit non-power-of-two cardinalities. The tested natural data simply does not contain enough of this structure for a useful win.

2. Scope Correction: Which “Attempt 4” Is This?

Original idea: this deck

Two adjacent scalar quantizers share one packed word. Their numbers of levels may be arbitrary positive integers.

Separate branch: not this result

The structured-2D method learns a reusable two-dimensional shape with affine transforms. It was evaluated separately and is currently set aside.

Why the distinction matters

Mixed-radix changes how a fixed number of bit patterns is divided between two scalar coordinates. It does **not** learn a shared 2D vector codebook.

3. The Original Intuition

Four bits give exactly $2^4 = 16$ possible labels for one coordinate pair.

Power-of-two allocation

Each coordinate receives 1, 2, 4, 8, ... levels. A typical balanced split is $4 \times 4 = 16$.

Mixed-radix allocation

Allow any integers K_1, K_2 as long as

$$K_1 K_2 \leq 16.$$

This includes $3 \times 5 = 15$.

Hypothesis

If one coordinate really needs 3 values and the other needs 5, mixed-radix can retain all 15 combinations. Separate binary rounding would need $4 \times 8 = 32$ states and cannot fit in four bits.

4. Running Example: Pack a 3×5 Grid

Let $z_1 \in \{0, 1, 2\}$ and $z_2 \in \{0, 1, 2, 3, 4\}$. Store

$$\text{label} = z_1 + 3z_2.$$

	$z_2 = 0$	1	2	3	4
$z_1 = 0$	0	3	6	9	12
$z_1 = 1$	1	4	7	10	13
$z_1 = 2$	2	5	8	11	14

Every real combination has a unique label; label 15 is simply unused.

What the power-of-two control must do

With only 16 labels it may choose 4×4 , but then five distinct values of the second coordinate must share only four reconstructed values. At least two real prototypes can receive the same stored code.

5. Training, Encoding, and Querying

Training, using base/index data only:

- for each adjacent coordinate pair:

 - estimate scalar reconstruction cost for each level count

 - choose K_1, K_2 minimizing total cost, subject to $K_1 * K_2 \leq 2^B$

Encoding one database vector:

- z_1 = nearest scalar centre on coordinate 1

- z_2 = nearest scalar centre on coordinate 2

- store label = $z_1 + K_1 * z_2$

Querying:

- build the complete 2^B -entry distance table for each pair

- decode one label per pair, add table values, keep the best 100 IDs

Controlled difference

From here on, A means arbitrary integer level counts and D means the power-of-two control. They use the same pairs, bytes, candidates, table builder, and scan loop. Only the allowed values of K_1, K_2 differ.

6. Closest Work Already Covers Most Primitives

Muresan–Effros	exact one-dimensional quantization with an arbitrary integer number of reconstruction levels
Brandt transform coding	PCA, bit allocation, trained scalar quantizers, packed words, and table-based ANN scanning
Adaptive Bit Allocation PQ	unequal integer bit budgets across independent PQ subspaces
Finite Scalar Quantization	Cartesian products of scalar level sets, including non-power-of-two factor sizes, mapped to one integer index
Quicker ADC	packed irregular-width PQ codes and the resulting lookup/SIMD costs
PQ / OPQ / block VQ	more general subspace codebooks or learned space decompositions under comparable payloads

Important boundary

No one paper is byte-for-byte identical to our implementation, but arbitrary level counts, bit allocation, mixed-radix indexing, packed lookup, and learned space decomposition are established primitives.

7. What Could Still Count as a Contribution?

Not enough by itself

A new combination of integer cardinalities, a lower reconstruction objective, or a different packing layout is vulnerable to the objection that it directly composes known quantization and lookup techniques.

The narrow remaining empirical claim

Under the same database payload and a fully counted query consumer, arbitrary cardinalities must produce a material, reproducible Recall–speed Pareto point that strong PQ/OPQ and power-of-two controls do not reach.

- Count codebooks, radix metadata, table construction, lookup work, index build time, and memory—not payload bytes alone.
- Show that the gain comes from arbitrary cardinalities, not a stronger transform, coordinate grouping, optimizer, or extra query work.
- The following experiments test this narrow claim; they do not claim that the mixed-radix formula itself is new.

8. What We Actually Compared

A: arbitrary

D: power-of-two

PQ / OPQ

coordinate view

data

payload

index partitions

quality

speed

search effort

arbitrary integer level counts, such as $15 \times 17 = 255$

identical method, but each level count must be a power of two

Product Quantization / Optimized Product Quantization baselines

first 128 PCA residual coordinates, forming 64 pairs

SIFT1M and GIST1M

exactly 32 or 64 code bytes per database vector

database divided into 1,024 or 4,096 coarse regions

Recall@100: how many of the true nearest 100 are returned

single-query latency and 12-core queries per second

nine fixed numbers of regions searched per query

Fairness rule

A and D search exactly the same candidates and perform the same table work. The allocation was fixed before reading query results.

9. What the Natural-data Allocator Chose

The 128-coordinate view forms 64 coordinate pairs; each pair contains two PCA coordinates and shares one packed label.

32-byte vectors: 4 bits per pair

All 64 coordinate pairs selected 4×4 .
Therefore A and D were literally the same representation and returned identical rankings.

64-byte vectors: 8 bits per pair

Depending on the dataset and index configuration, only 10–24 of the 64 coordinate pairs differed from 16×16 .
Almost all were 15×17 or 17×15 , using 255 of 256 paid states.

Early mechanism diagnosis

The reconstruction-based allocator found no large unused-cardinality problem on SIFT/GIST. At 64 bytes it mostly moved only one level from one coordinate to its partner.

10. Natural-data Recall Result

At the same number of searched database regions, A-minus-D Recall@100 was:

Dataset	total index partitions	minimum	maximum
GIST1M	1,024	-0.00096	+0.00054
GIST1M	4,096	-0.00106	+0.00017
SIFT1M	1,024	-0.000097	+0.000066
SIFT1M	4,096	-0.000033	+0.000436

Meaning in plain language

The largest gain was 0.054 percentage points; the largest loss was 0.106 percentage points. The sign changed across datasets, index sizes, and search ranges, so there is no stable quality improvement.

11. Speed and Baseline Result

- A/D throughput ratio: 0.9918 to 1.0069; configuration medians are essentially 1.
- A/D p95 latency ratio: 0.9897 to 1.0251.
- This is expected: A and D deliberately execute the same complete-table query routine and nearly the same representation.
- PQ and OPQ generally reach higher Recall in the useful high-quality region.

Natural-data decision

The fixed-adjacent mixed-radix method provided neither clearly higher Recall at similar speed nor clearly higher speed at similar Recall. A tiny nonzero difference is not a database-systems contribution.

Complete matrix: 512 accepted passes; 32.36 CPU-hours; 17.09 wall-hours; all eight correctness tests pass.

12. Why Add a Synthetic Positive Control?

The natural result left two different explanations:

- 1 the mixed-radix mechanism is valid, but the required data structure is absent or weak; or
- 2 the implementation cannot transmit a radix advantage into query ranking, even when such an advantage exists.

Decisive construction

Create exact Cartesian prototypes with the cardinalities required by the intuition. Put 100 identical database vectors at every prototype, place one query exactly on each prototype, search every candidate, and ask whether all 100 true duplicates are recovered.

Positive cases are 3×5 in 4 bits and 15×17 in 8 bits. Matching 4×4 and 16×16 cases are zero-effect controls.

13. Synthetic Result: The Mechanism Reaches Recall

Case	A shape	D shape	merged prototype pairs A/D	Recall A/D
positive B4	3×5	4×4	0 / 3	1.000 / 0.800
null B4	4×4	4×4	0 / 0	1.000 / 1.000
positive B8	15×17	16×16	0 / 15	1.000 / 0.941
null B8	16×16	16×16	0 / 0	1.000 / 1.000

Validity checks

In both null controls A and D produced identical rankings. In both positive cases the rankings differed exactly where D merged prototypes. Two complete runs were byte-for-byte identical.

14. Put the Two Results Together

Question	Natural SIFT/GIST	Synthetic witness
Can A choose non-dyadic radices?	rarely, and only weakly at 64 bytes	yes, exactly 3×5 and 15×17
Does this avoid collisions?	no material measured benefit	yes, 3 and 15 control collisions removed
Does it improve Recall?	tiny, mixed-sign changes	yes, by 0.200 and 0.0588
Does it beat PQ/OPQ?	no better Recall-speed trade-off	not tested or claimed

Combined conclusion

The mechanism is real but its necessary structure is not sufficiently strong in the tested natural representation. The synthetic result explains the negative result; it does not reverse it.

15. Coordinate Regrouping Was Not the Missing Contribution

Allowing non-adjacent coordinate pairs improved Recall substantially over the fixed pairing $(0, 1), (2, 3), \dots$. We then compared our minimum-cost perfect matching with our two-coordinate specialization of the Eigenvalue Allocation rule in parametric OPQ (Ge et al., Section 3.2.4).

	search 64 cells	search 1,024 cells
SIFT1M: matching minus 2D allocation control	+0.000253	+0.000349
GIST1M: matching minus 2D allocation control	+0.000320	-0.000160

Interpretation

Under our frozen two-coordinate scalar consumer, “do not pair coordinates mechanically” is useful, but the tested OPQ-P-inspired control already explains nearly all of the gain. This is not a comparison with full OPQ. Arbitrary integer radices still add at most about 0.00057 Recall.

16. Why Reconstruction Error Was Not Enough

Recall depends on the ordering of complete estimated distances, not on the average reconstruction error of one coordinate block. Here a segment means a block of coordinates encoded and scored together.

The diagnostic question

For each database vector being re-encoded, replace one block's approximate distance by its true distance while leaving every other block approximate. How many wrongly ordered vector pairs become correct?

- This is an **idealized exact-distance control**: it uses information that a deployable compressed code does not store.
- A single coordinate block repaired only about 43–44% of ordering errors and missed the frozen absolute-improvement rule.
- Jointly replacing a 192-coordinate block at 6 bits per coordinate and a 320-coordinate block at 4 bits per coordinate repaired **71.1–75.6%**.

What this did and did not establish

The two blocks jointly contain many potentially removable ordering errors. True-distance replacement is not an encoder; the Recall ordering test can make separate errors look jointly important when they cancel each other.

17. Can Base Data Learn a Joint Encoder?

With the existing SAQ index layout fixed, every legal code for the first block can be combined with every legal code for the second block. Standard reconstruction and worst-case errors still split into one objective per block; joint benefit can only come from error correlations in chosen test directions.

Illustrative cancellation

On training directions, two block errors might be $(+2)$ and (-2) , so the total error is zero. On different directions they might become $(+2)$ and $(+2)$, so the same jointly selected pair has total error four. A cancellation seen during training need not survive new directions.

Falsifiable hypothesis

Choose the two codes jointly using directions constructed from base vectors, not benchmark queries. Freeze the choices and test them on a disjoint half. Require at least 5% lower mean squared estimator error than separate selection in both train/test swaps, and improvement on at least 60% of database vectors.

18. Train on One Half, Test on the Other, Then Swap

No benchmark query or ground truth was read. Each listed database vector is one whose two stored block codes are being reconsidered.

```
for each of 32 fixed database vectors:
  alternatives = stored code + every legal one-coordinate +/-1 move
  train on base-derived directions 0..127:
    SEPARATE = best code for each block separately
    JOINT     = best of every possible two-block code combination
  freeze both choices; test on directions 128..255
repeat after swapping the two direction halves
```

Both methods use the same alternatives, bytes, scaling calculation, and distance formula. The post-hoc test-set ceiling chooses only after seeing the test directions; it shows remaining local opportunity but is not a deployable method.

19. Joint Selection Did Not Generalize

	train half 0	train half 1
mean squared error reduction vs separate	+0.0805%	-1.3644%
database vectors improved	15/32	13/32
joint and separate choices differ	31/32	29/32
post-hoc test-set ceiling	16.83%	15.92%

The first row is error reduction, so larger is better; the negative value means that joint selection made the error worse.

Decision: the joint rule failed on the separate test half

Good code pairs exist after seeing the test directions, but choices learned from the other half do not predict them. The apparent opportunity comes from direction-specific error cancellation, not a stable base-trained rule.

20. Final Claim Boundary

What remains valid

Mixed-radix packing can preserve combinations of non-power-of-two level counts and prevent collisions in data deliberately having that structure.

What the natural-data evidence rejects

Neither fixed adjacent mixed-radix, optimized coordinate pairing, nor the tested base-trained two-block joint objective produces a stable new method with a better Recall–speed trade-off than the established comparisons.

Research decision

Close this line. Do not enlarge the local candidate set, change the data split or threshold, or inspect benchmark queries to rescue it. Continuing would require a genuinely new question—for example changing the distance formula or storing information shared across coordinate blocks—followed by a fresh review of the closest prior work and its costs.

References and Evidence

- Muresan and Effros. "Quantization as Histogram Segmentation." DCC 2002.
- Brandt. "Transform Coding for Fast Approximate Nearest Neighbor Search." CVPR 2010.
- Jégou, Douze, and Schmid. "Product Quantization for Nearest Neighbor Search." TPAMI 2011.
- Ge et al. "Optimized Product Quantization." CVPR 2013, Section 3.2.4, "Eigenvalue Allocation."
- Guo et al. "Adaptive Bit Allocation Product Quantization." Neurocomputing 2016.
- André et al. "Quicker ADC." TPAMI 2021.
- Mentzer et al. "Finite Scalar Quantization." ICLR 2024.
- Gao and Long. "RaBitQ." SIGMOD 2024.
- Li et al. "SAQ." 2025.
- Frozen mixed-radix design:
mixed_radix_query_design_2026_07_30.md.
- Natural-query result:
mixed_radix_query_results_2026_08_01.md.
- Synthetic mechanism result:
mixed_radix_synthetic_witness_result_2026_08_01.md.
- Closest grouping baseline:
nonadjacent_pairing_closest_baseline_result_2026_08_03.md.
- Exact component and joint-component oracle results:
saq_component_oracle_diagnostic_result_2026_08_03.md;
saq_joint_component_oracle_result_2026_08_03.md.
- Static objective audit and base-trained cross-fit result:
saq_two_segment_feasible_objective_static_audit_2026_08_03.md;
saq_joint_objective_viability_result_2026_08_04.md.